

# Inside VOLT: Designing an Open-Source GPU Compiler

Shinnung Jeong<sup>1</sup>, Chihyo Ahn<sup>1</sup>, Huanzhi Pu<sup>1</sup>, Jisheng Zhao<sup>1</sup>,  
Hyesoon Kim<sup>1</sup>, Blaise Tine<sup>2</sup>

<sup>1</sup>Georgia Institute of Technology, <sup>2</sup>University of California, Los Angeles



# GPUs Power Modern Computing - Openness Remains Limited



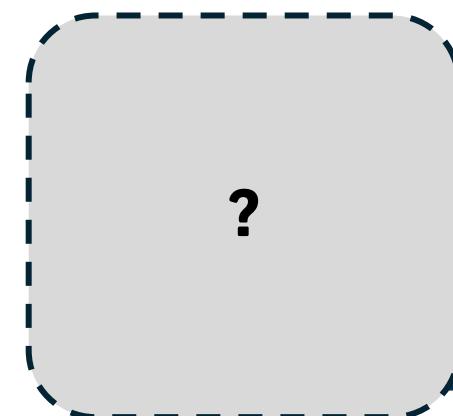
Core GPU microarchitecture and compiler internals  
remain largely inaccessible to researchers

# What Research Needs Beyond Exposed APIs



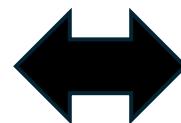
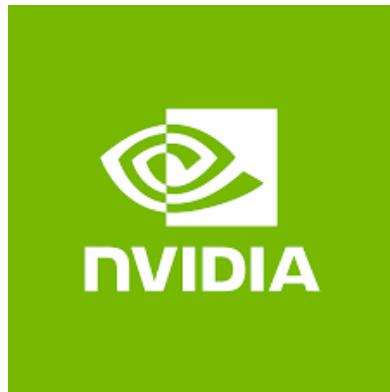
AMD

intel



GPU research would benefit from openness beyond exposed APIs

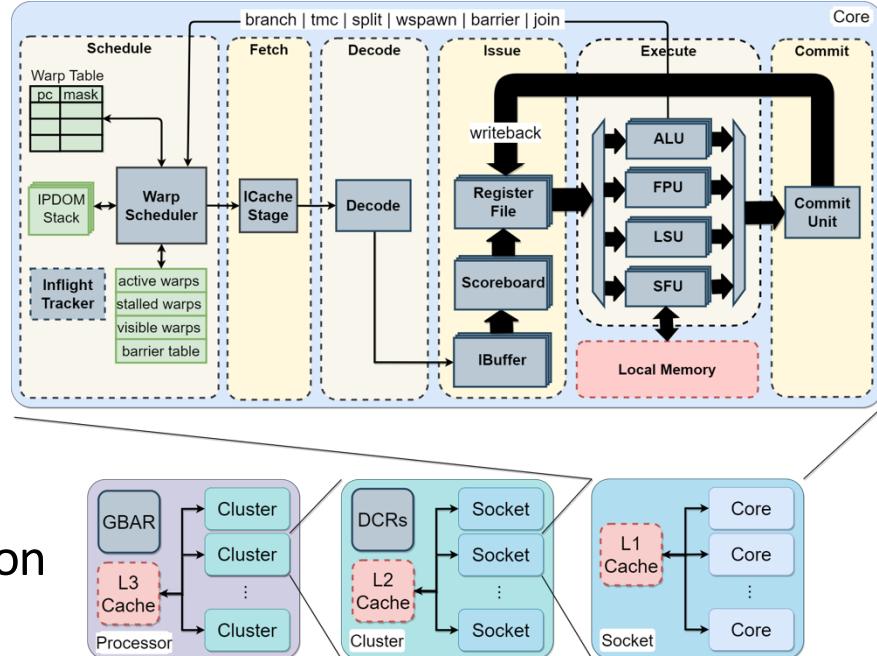
# What Research Needs Beyond Exposed APIs



Vortex GPU provides a fully open-source GPU platform

# Vortex GPU

- RISC-V based GPGPU
- Highly reconfigurable hierarchical architecture
- SIMT execution model
- Explicit hardware support for SIMT behavior, including control-flow divergence and barriers
- Drivers for host-device communication



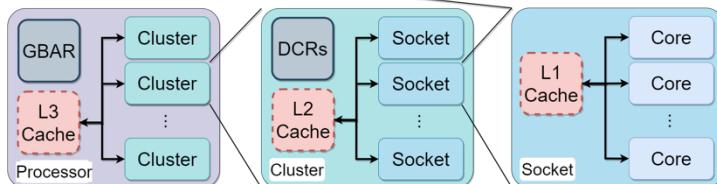
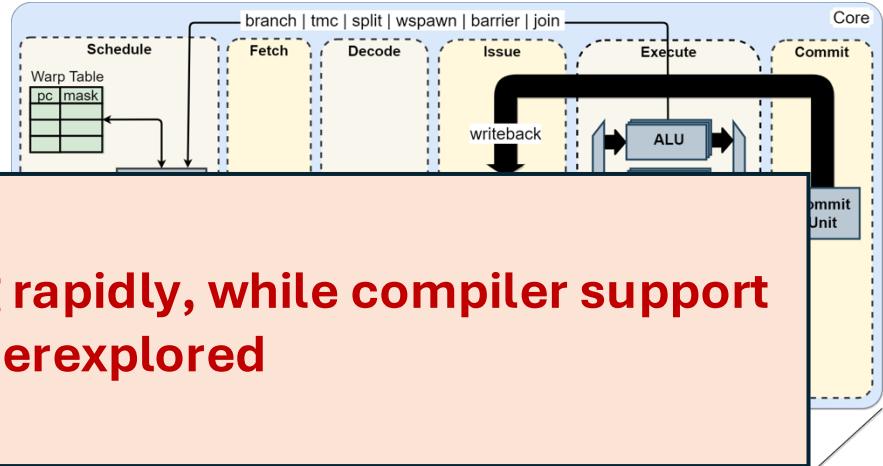
# Vortex GPU

- RISC-V based GPGPU
- Highly reconfigurable hierarchical

**Open GPU hardware is advancing rapidly, while compiler support remains underexplored**

divergence and barriers

- Drivers for host-device communication



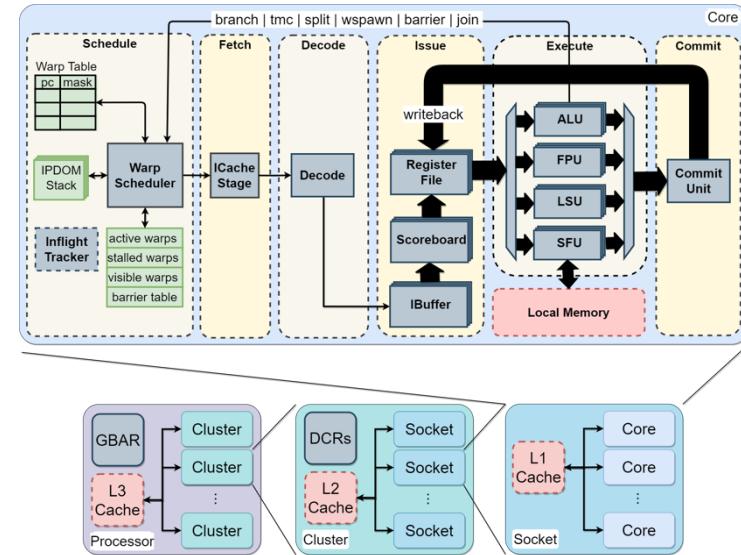
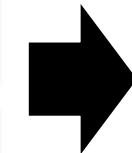
# Program Execution on Vortex GPU

```
__kernel void foo(){  
    tid = get_global_id(0);  
    ...  
    barrier(...);  
}
```

```
__host int main(){  
    ...  
    cl_mem A_clmem = clCreateBuffer(context, ...);  
    clEnqueueNDRangeKernel(command_queue, ...);  
    ...  
}
```

## Program

- Kernel code with control flow & special func
- Host Code with device communication



## Vortex GPU

- Kernel binary using the Vortex GPU ISA
- Host executable using Vortex driver

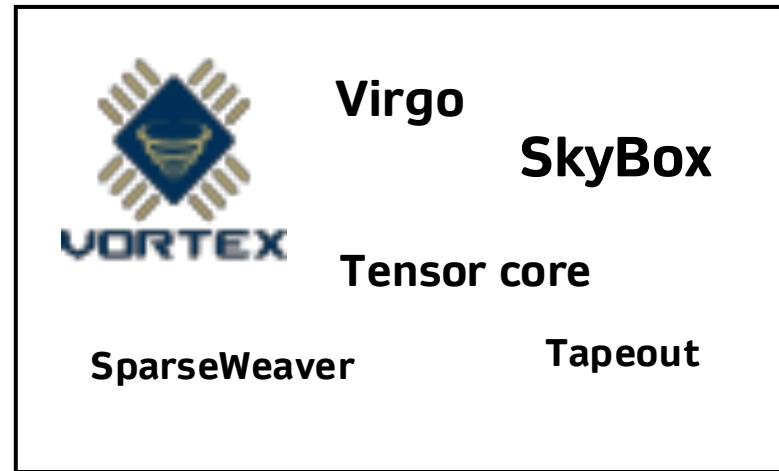
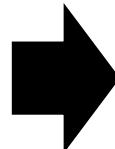
# Growing GPU Variants on Vortex

```
__kernel void foo(){  
    tid = get_global_id(0);  
    ...  
    barrier(...);  
}
```

```
__host int main(){  
    ...  
    cl_mem A_clmem = clCreateBuffer(context, ...);  
    clEnqueueNDRangeKernel(command_queue, ...);  
    ...  
}
```

## Program

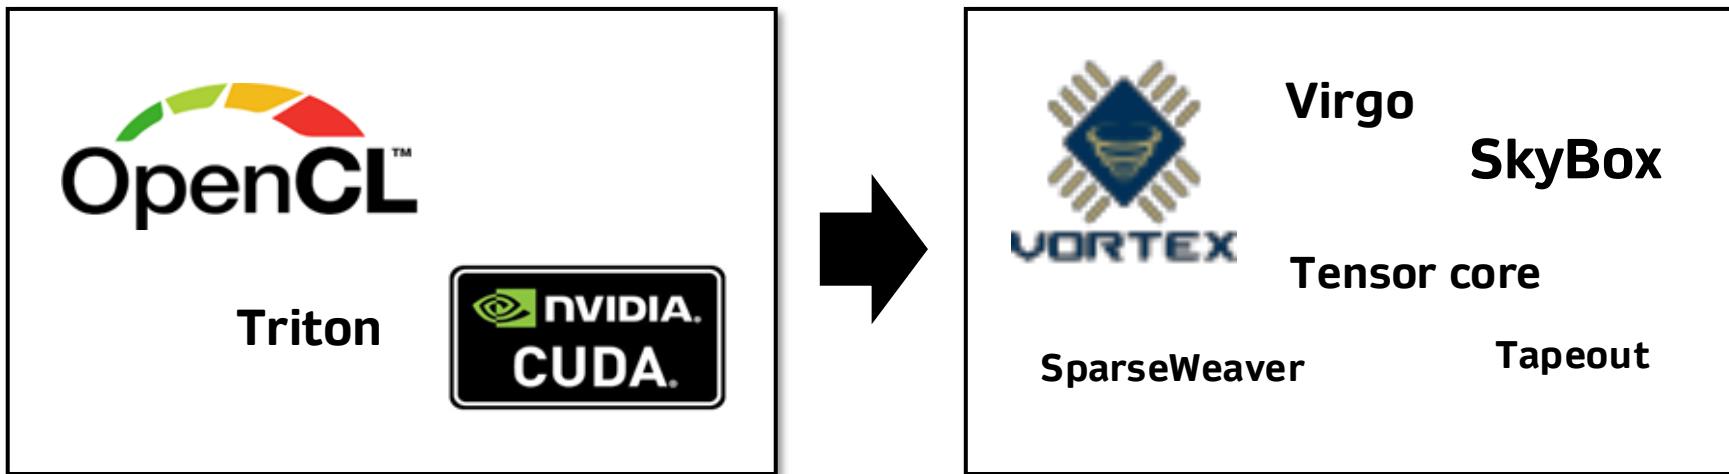
- Kernel code with control flow & special func
- Host Code with device communication



## Vortex GPU Variant

- Kernel binary using the **Extended** Vortex GPU ISA
- Host executable using **Extended** Vortex driver

# Programs from Diverse Frontends



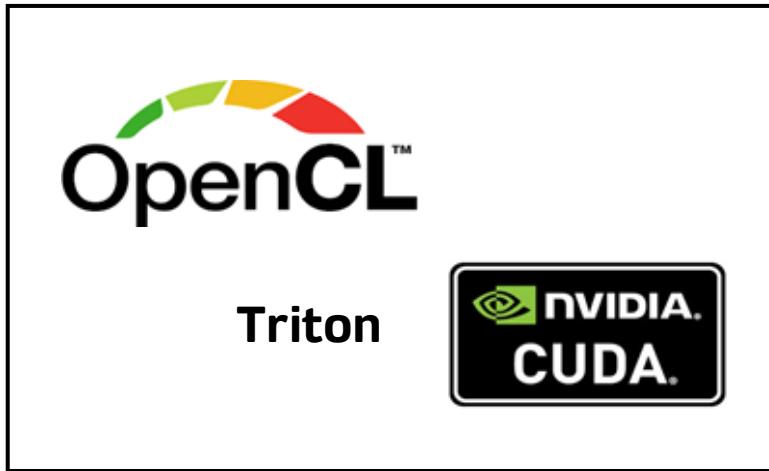
## Program with Diverse Frontend

- Kernel code with **Frontend-specific semantics**
- Host Code with **Frontend-specific runtimes**

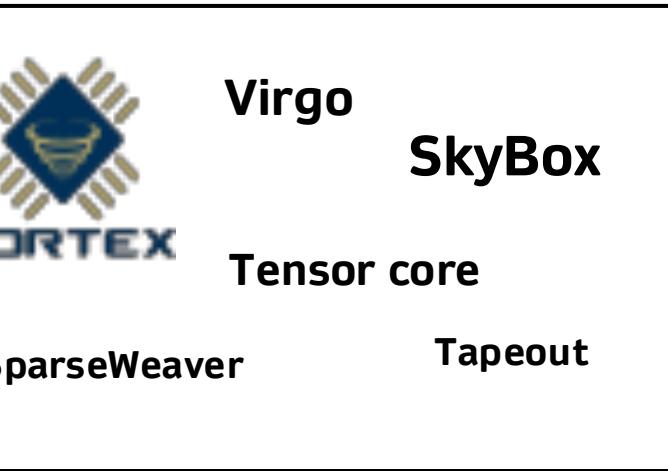
## Vortex GPU Variant

- Kernel binary using the **Extended** Vortex GPU ISA
- Host executable using **Extended** Vortex driver

# Managing Growing GPU Variants and Diverse Frontends



Program with Diverse Frontend



Vortex GPU Variant

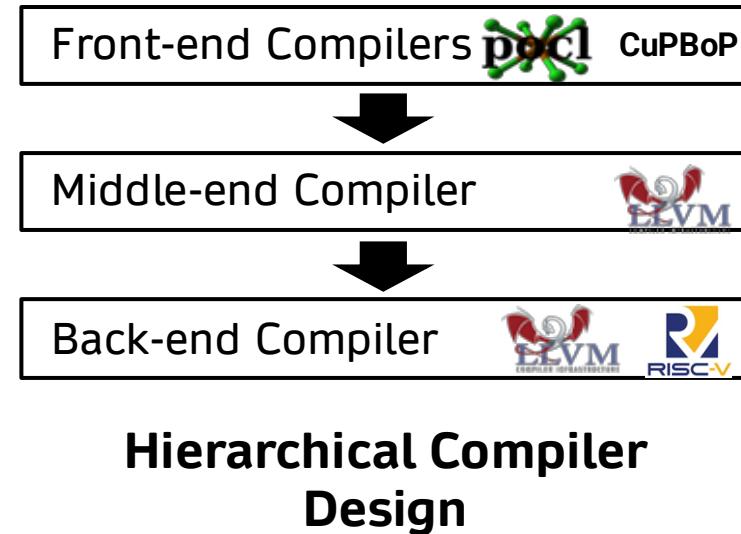
**This Calls for a Well-Designed Compiler Framework**

# Key Challenges in Designing VOLT

- **Challenge 1: Make SIMT-aware code generation and optimizations portable**
  - How can we preserve SIMT-aware code generation and optimizations while remaining portable across Vortex variants and potentially other open GPUs?
- **Challenge 2: Support multiple frontends with minimal maintenance overhead**
  - How can we design the compiler so that new frontends and GPU variants can be added with minimal refactoring?

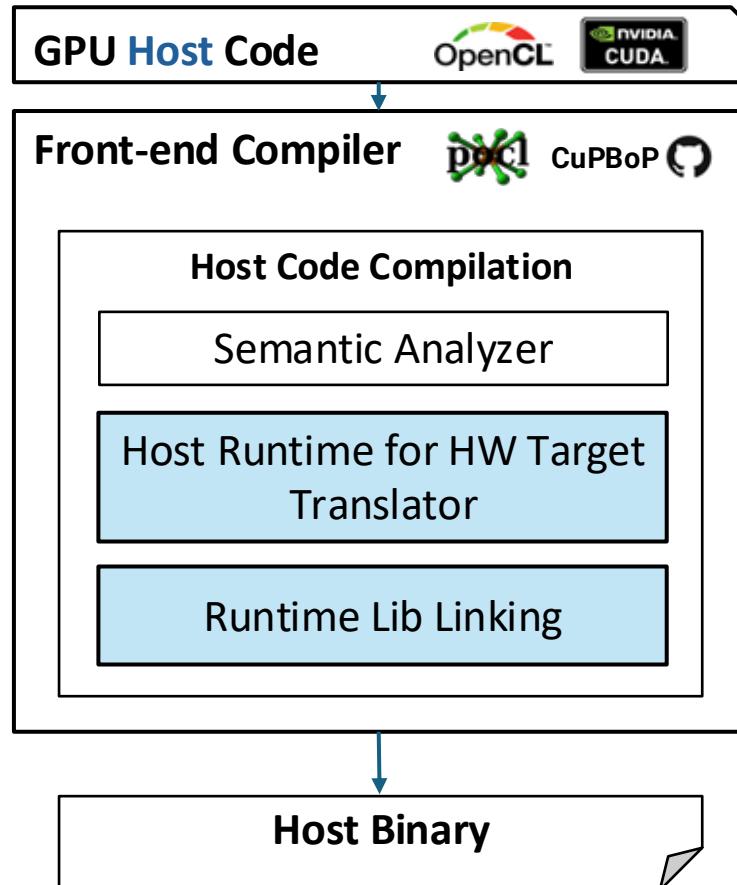
# Vortex-Optimized Lightweight Toolchain (VOLT)

- Design Choice 1
  - Hierarchical compiler design
  - Centralize fundamental SIMT-related analyses and optimizations in the middle-end
- Design Choice 2
  - Leverage existing open-source components as much as possible



# Front-end Compiler

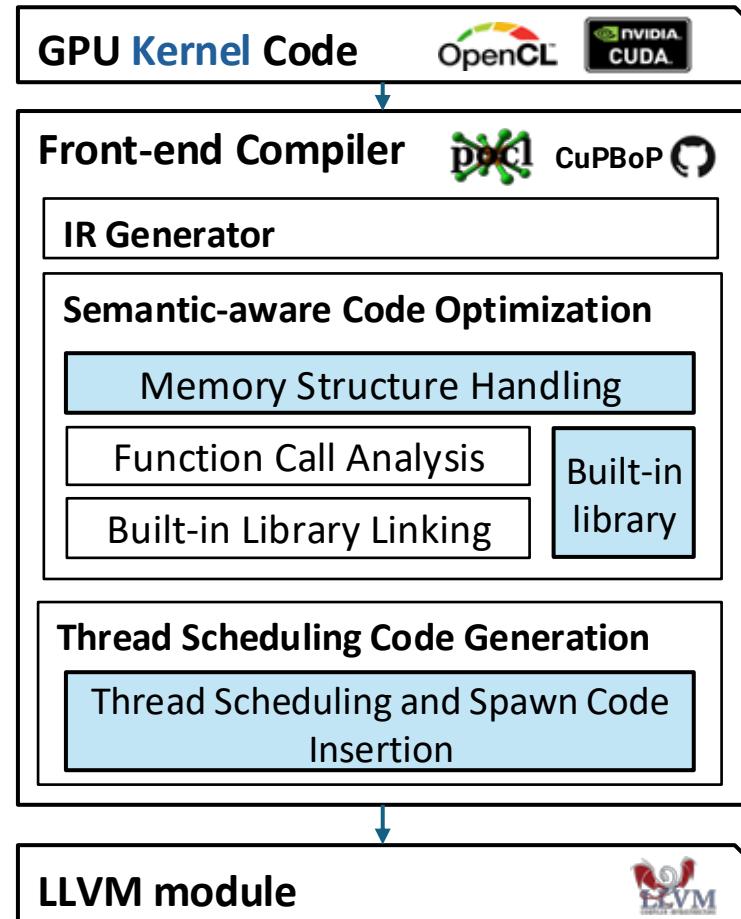
- Handles Both Host and Kernel Code
- **Host Code**
  - Compiled for the host hardware
  - Translates the frontend host functions using the target runtime library



Blue boxes indicate components extended or implemented for the Vortex GPU

# Front-end Compiler

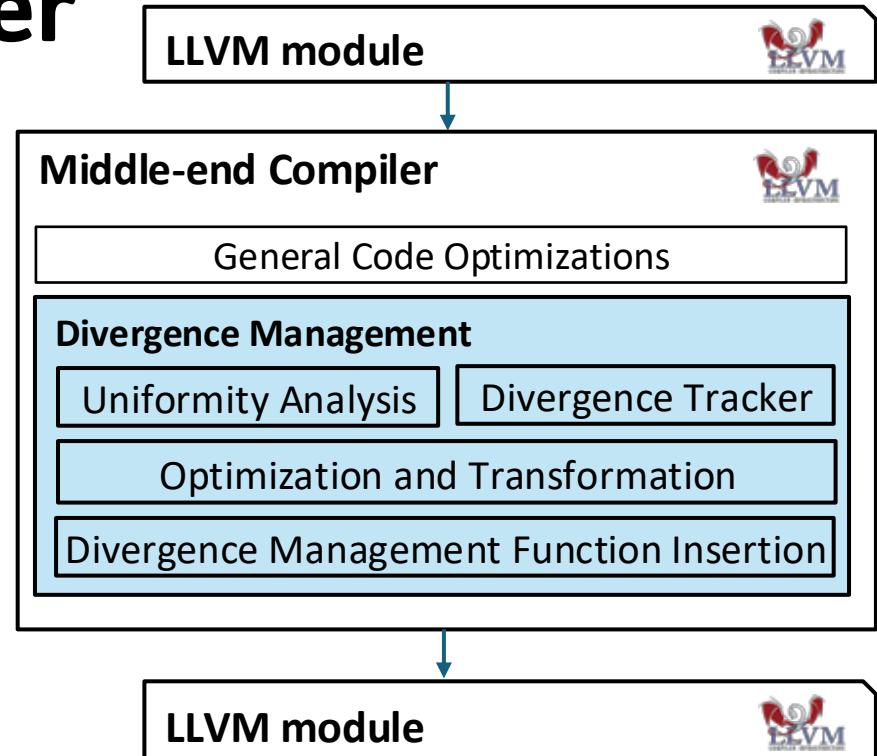
- Handles Both Host and Kernel Code
- **Kernel Code**
  - Performs semantics-aware code optimizations
  - Transforms special functions using built-in libraries
  - Inserts thread scheduling and spawning code for device initialization



Blue boxes indicate components extended or implemented for the Vortex GPU

# Middle-end Compiler

- Target-independent optimizations
- Minimal target-specific logic
- Thread Divergence management treated as a first-class concern



Blue boxes indicate components extended or implemented for the Vortex GPU

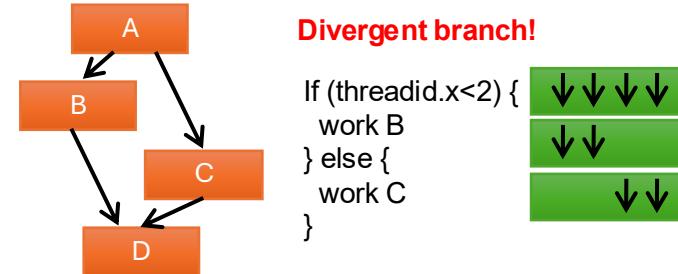
# Thread Divergence Management

- SIMD

- Single instruction applied to multiple data elements
- Uniform control flow
- Divergent branches are serialized

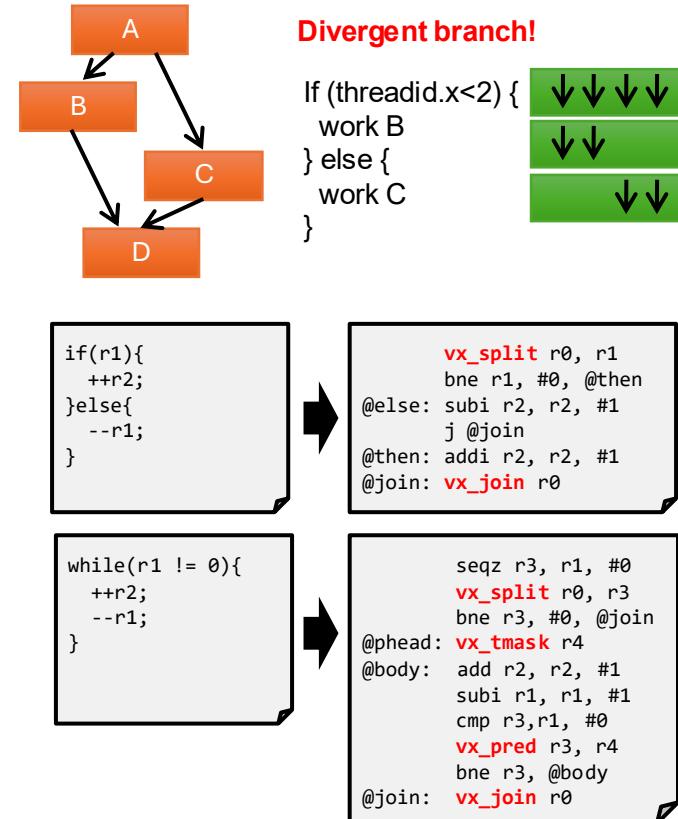
- SIMT

- Threads execute in lockstep under a shared instruction stream
- Each thread has its own registers and stack
- Control-flow divergence is handled via masking and serialization



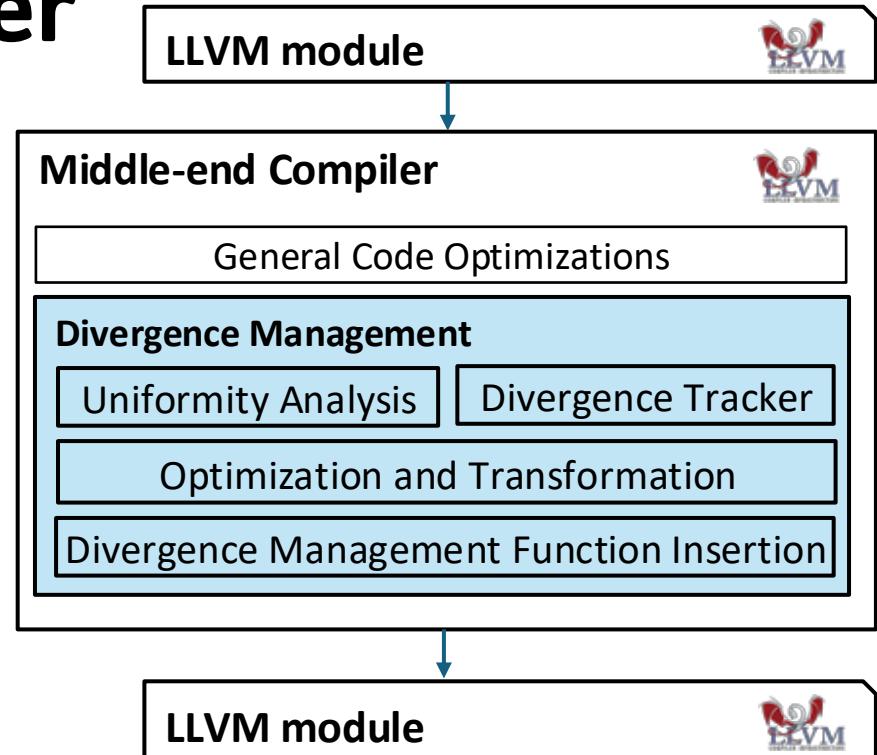
# Thread Divergence Management

- Divergence management microarchitecture
  - IPDOM stack for handling divergence and reconvergence
  - Per-warp thread mask for control-flow management
- ISA support
  - Thread mask control instruction
  - Split and join instructions for If
  - Predicate instruction for loop



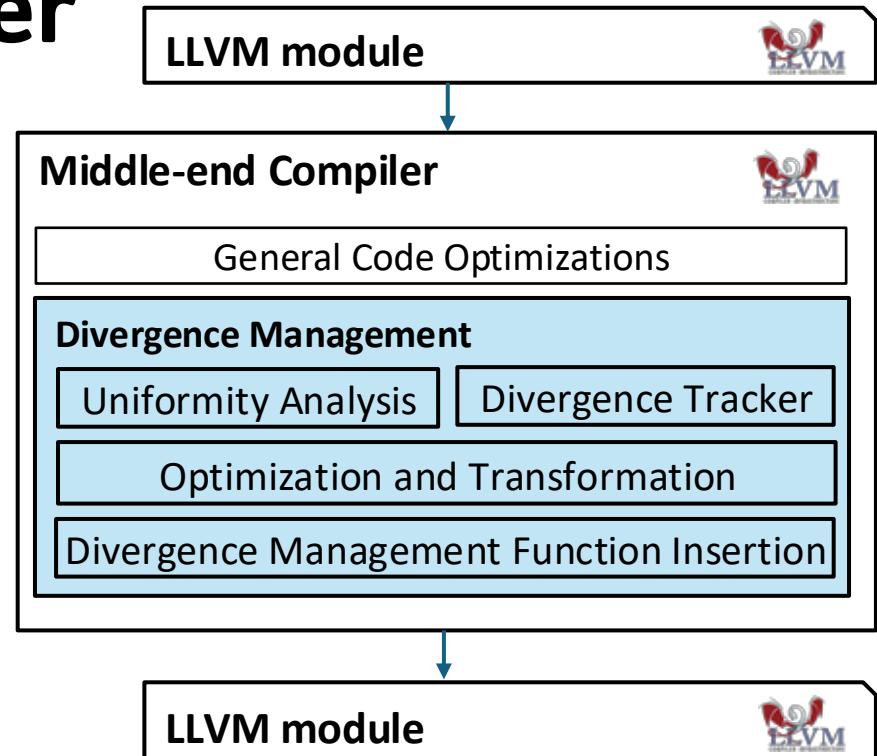
# Middle-end Compiler

- **Uniformity Analysis**
  - Find divergent instructions and always uniform instructions
- **Divergence Tracker**
  - Mark Function arg, returns, Atomic as Divergent
  - Mark Control and Status Register info as Uniform
  - Propagate information



# Middle-end Compiler

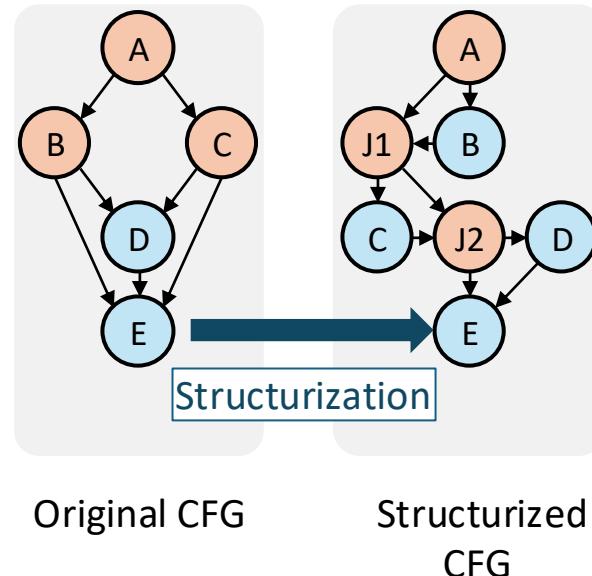
- **Uniformity Analysis**
  - Identify divergent instructions and always-uniform instructions
- **Annotation analysis**
  - Leverage annotation information from the user or the front-end compiler
- **Function argument analysis**
  - Determine the uniformity of each function's arguments



Blue boxes indicate components extended or implemented for the Vortex GPU

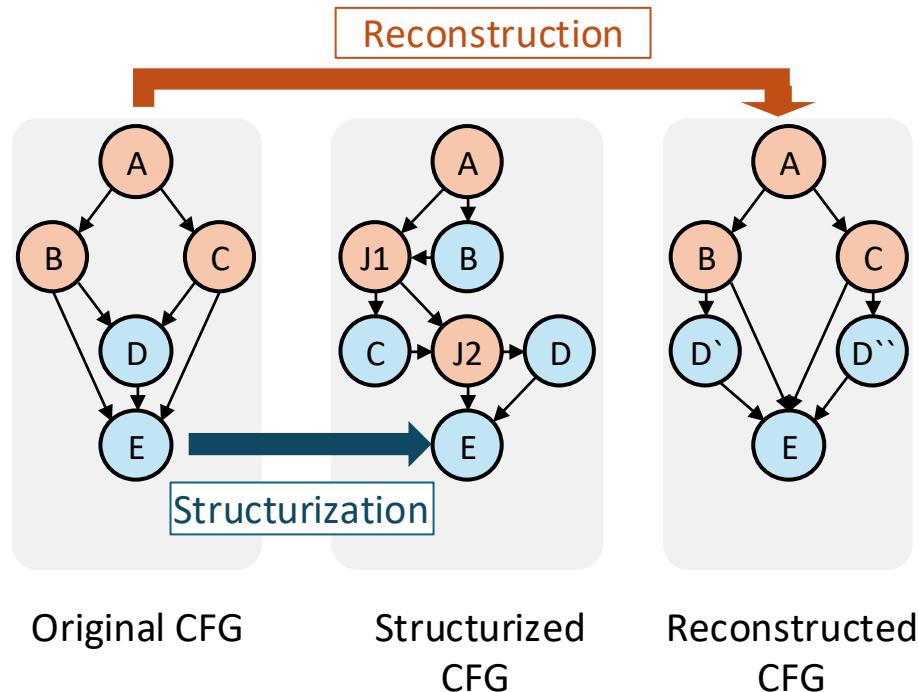
# Middle-end Compiler

- Divergence Optimization
  - Code Simplification
    - Simplifies control flow
    - Canonicalizes control flow to simplify subsequent divergence handling
  - Control-Flow Structurization
    - Transforms the CFG into a structured form to provide stable join points
    - Essential for IPDOM-based divergence and reconvergence handling



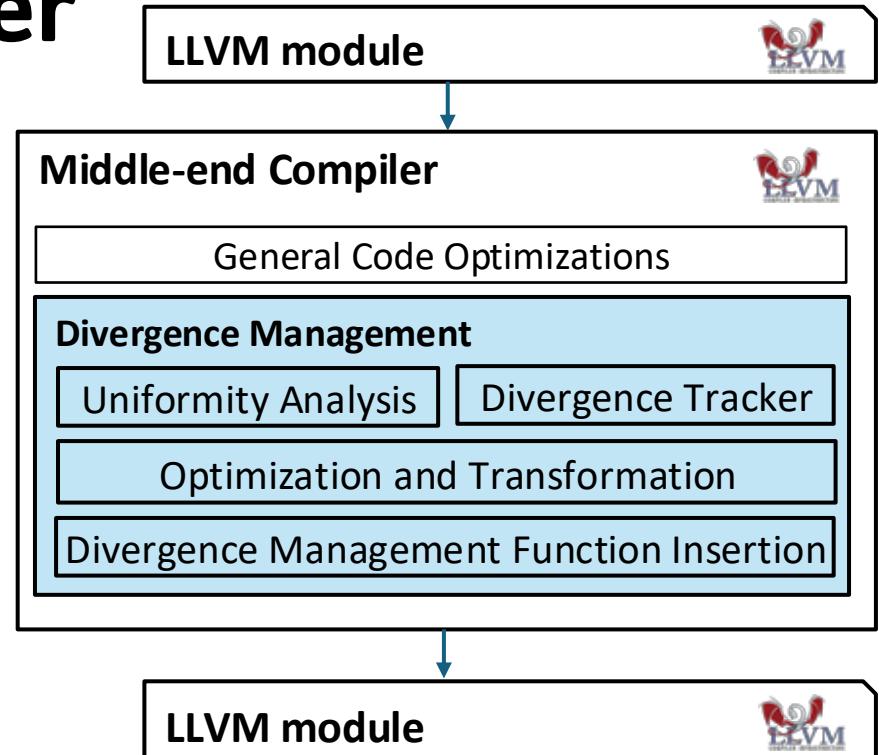
# Middle-end Compiler

- Divergence Optimization
  - Control-Flow Reconstruction
    - Selectively duplicates basic blocks to reduce predicate computation
  - Divergence Operation Lowering
    - Normalizes divergence operations by rewriting them into equivalent branch-based control flow



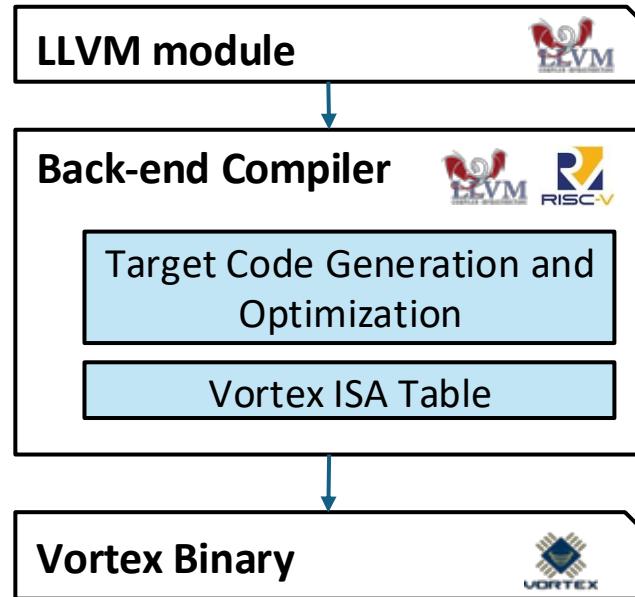
# Middle-end Compiler

- Divergence Intrinsic Insertion
  - Split and Join Insertion
    - Handling divergent branch with `vx_split` and `vx_join` instruction
  - Loop predicate Insertion
    - Handling divergent loops with `vx_pred`



# Back-end Compiler

- Handles target-specific optimizations and final code generation
- Built on the RISC-V target compiler
- Extends the ISA table with Vortex-specific ISA extensions
- Generates Vortex target binaries



Blue boxes indicate components extended or implemented for the Vortex GPU

# VOLT v1.0 Release



- Our tool is available in our Git repository
  - <https://github.com/vortexgpgpu/Volt/releases/tag/v1.0>
- The first release supports CUDA 12.1 and OpenCL 3.0
- Provides an end-to-end PoCL / CuPBoP / LLVM compilation pipeline targeting the Vortex GPU

A screenshot of the GitHub repository page for 'Volt'. The repository is public and has one branch ('master') and one tag ('v1.0'). The repository owner is 'shin0403' and they made a 'minor update'. The repository contains several sub-directories: 'cupbop', 'docs', 'llvm', 'pcl', 'tests', and 'vortex'. Each directory has a corresponding commit hash listed next to it.

Directory	Commit Hash
cupbop	@ 799891e
docs	
llvm	@ bb85f7e
pcl	@ f613f42
tests	
vortex	@ a4b6543

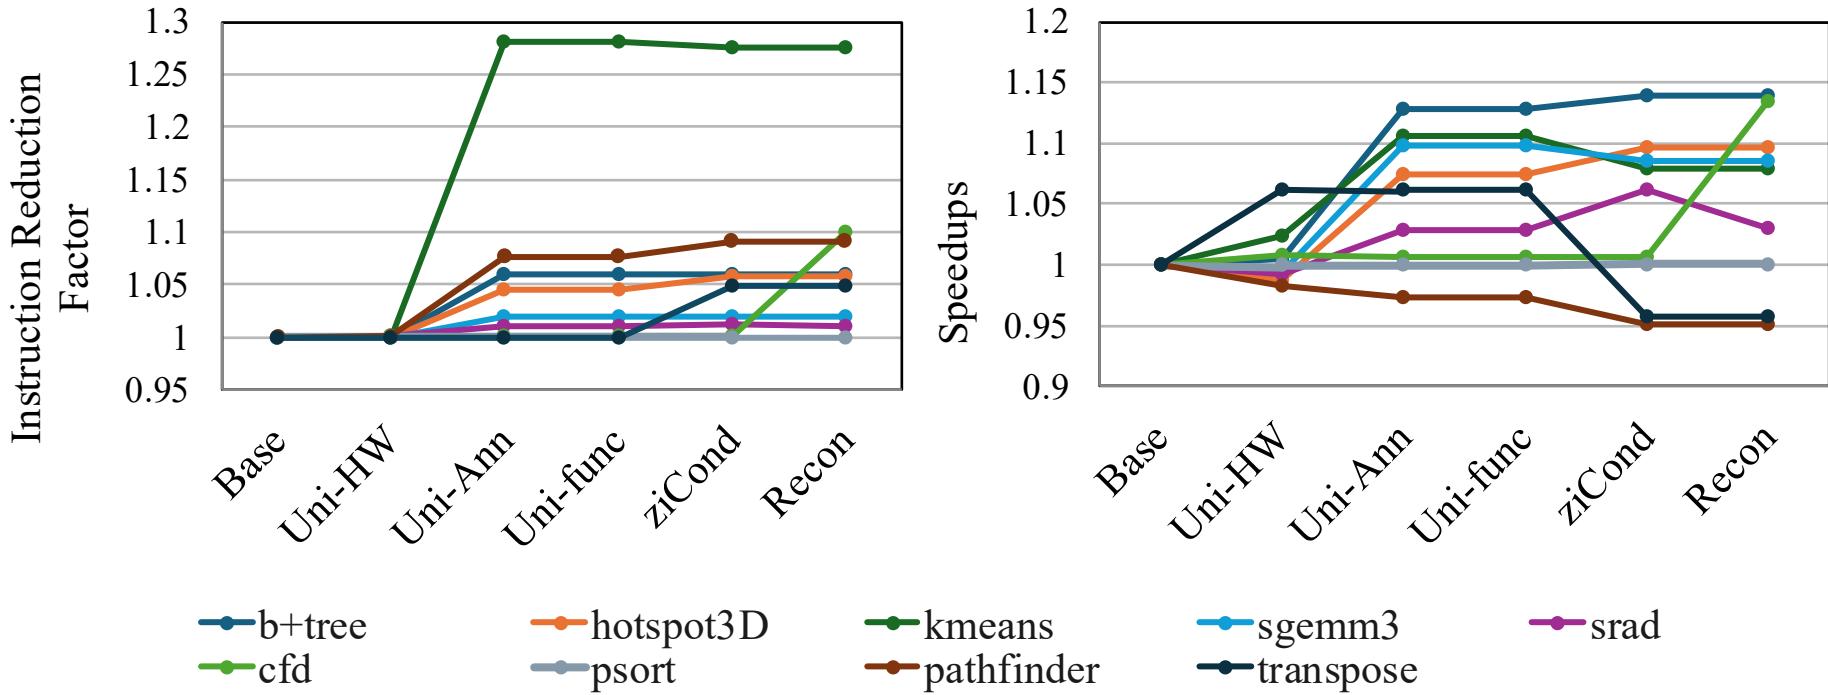
# Guide for VOLT Extensions



- Tutorial 1: Extending kernel functions
- Tutorial 2: Extending host functions
- Tutorial 3: Extending memory support with the Vortex memory hierarchy
  
- All tutorials include documentation under /Docs and runnable examples under /Tests.

0.index.md
1.getting_started.md
2.overview.md
3.PoCL_vortex.md
4.CuPBoP_vortex.md
5.LLVM_vortex.md
6.analysis_and_transform_passes.md
7.1.tutorial_kernel_extension.md
7.2.tutorial_host_extension.md
7.3.tutorial_shared_memory.md

# Divergence Management Optimization



# Research Directions with Volt & Vortex

## 1. Micro-architectural Compiler Optimization

- Fine-grained instruction & warp scheduling
- Latency hiding via instruction prefetching
- Adaptive workload distribution and occupancy control

## 2. Architecture-aware Optimization & Analysis

- Exploitation Heterogeneous core
- Explore reconfigurable options

## 3. HW / SW Co-design Exploration

- Co-design of custom GPU extensions and compiler support
- ISA or micro-architectural features
- Feedback-driven optimization between hardware, compiler, and runtime

# Conclusion

- **VOLT is a lightweight, extensible compiler toolchain optimized for Vortex**
- **Hierarchical design, Centralize SIMD-aware optimizations in the middle-end**
  - Enables portability across Vortex variants
- **Thread divergence management is treated as a first-class concern**
  - Uniformity analysis, divergence tracking, and control-flow optimizations
  - Improves performance and reduces instruction count
- **Extensibility and reproducibility**
  - Clear extension points for kernel, host, and memory support
  - Tutorials, documentation, and shared scripts enable easy adoption and extension



# Thank you

